

In this lecture, we will discuss...

- ✧ Rationale behind NoSQL
- ✧ Scaling Issues in RDBMS
- ✧ NoSQL: What is it?

Why RDBMS

- ✧ Relational Databases – popular and commonly used
- ✧ Initially designed for non distributed
- ✧ Low Cost RDBMS alternatives (PostgreSQL, MySQL, SQLite)
- ✧ Very Transactional - across tables and commands, and can even be transactional across distributed resources (XA) -- at a cost
- ✧ Supports Joins -- across multiple tables allowing for normalized forms of data to be stored once



Why NoSQL

- ✧ Explosion in data
- ✧ Object/Relational Impedance mismatch
 - Objects are **constantly being moved** in/out of tables/rows
- ✧ RDBMS normalization and joins are **powerful**, but add up in **cost**
 - Complex objects stored across many tables and rows can be **expensive** to handle



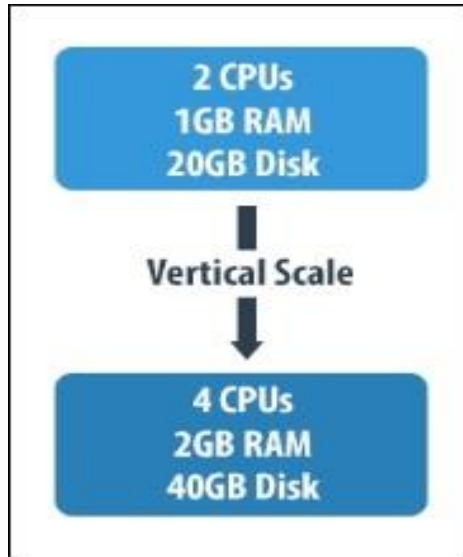
Why NoSQL

- ✧ “Big” data handling with better performance
- ✧ Supports unstructured data
 - Unique data type extensions can be easily integrated into existing collections
- ✧ Operational issues (scale, performance and availability)

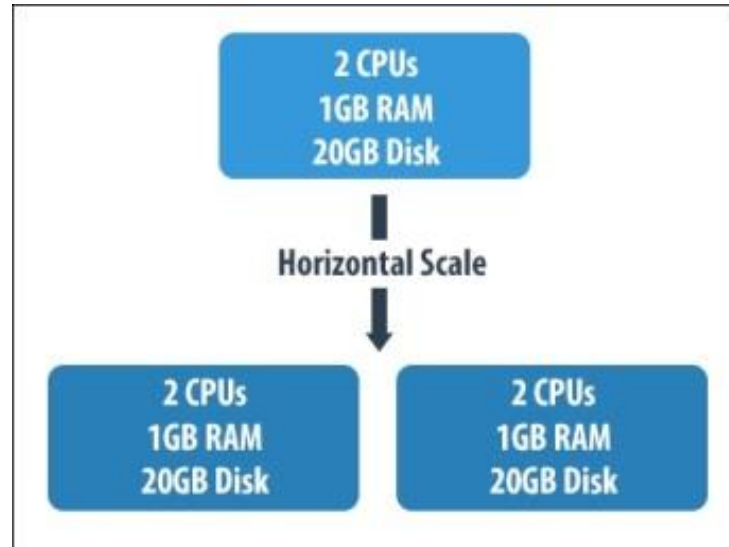


Scaling Out

Vertical Scaling



Horizontal Scaling



What is NoSQL

- ✧ Stands for “Not Only SQL”
- ✧ No Fixed Schema
- ✧ Non-relational data storage systems



Summary

- ✧ NoSQL – very popular and major companies especially social networking sites such as Twitter, Facebook, LinkedIn, and Digg use NoSQL DB
- ✧ Excellent performance and stability, fast and scalable and fairly simple model
- ✧ Supports unstructured format, which makes it very agile
- ✧ NoSQL is mostly gained when access patterns to complex objects are understood and modeled correctly up front



What's Next?

Categories of NoSQL

